



Hashing

Yi-Shin Chen

Institute of Information Systems and Applications

Department of Computer Science

National Tsing Hua University

yishin@gmail.com

Improve Retrieval Efficiency

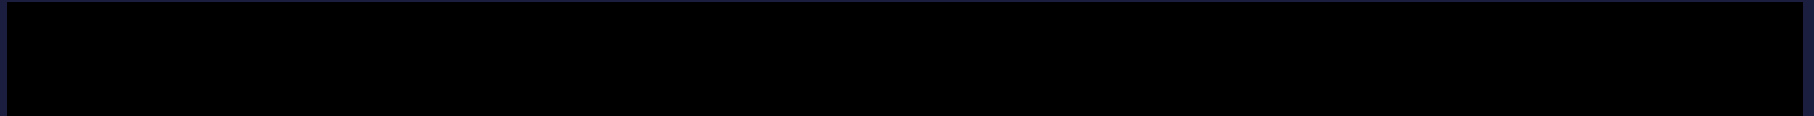
- Sort the list in a specific order before searching
- Approaches
 - Sort after a batch of insertion
 - Insertion time should be small
 - Chance of retrieval is rare
 - Insertion based on some sorting policy
 - Retrieval time should be small

Indexing

- Balanced (binary) search tree
 - Get, Insert and Delete take $O(\log n)$
- Hashing
 - Get, Insert and Delete take $O(1)$
 - Static hashing
 - Dynamic hashing




Static Hashing



Overview of Hashing

- The file blocks are divided into M equal-sized *buckets*
- The record with hash key value K is stored in bucket I
 - $i = h(K)$, and h is the *hashing function*




Hash Table

- Hash table (ht)
 - A container stores dictionary pairs
- Hash table is partitioned into b buckets
 - $ht[0], ht[1], \dots, ht[b-1]$
 - Each bucket holds s dictionary pairs (slots)
 - Usually $s=1$ which means each bucket can hold exactly one pair

0
1
N-1

Hash Function

- The *hash (address)* of the pair is determined by a **hash function**, $h(k)$
 - Hash function maps keys into buckets by returning an integer in the range 0 through $b-1$



Definitions

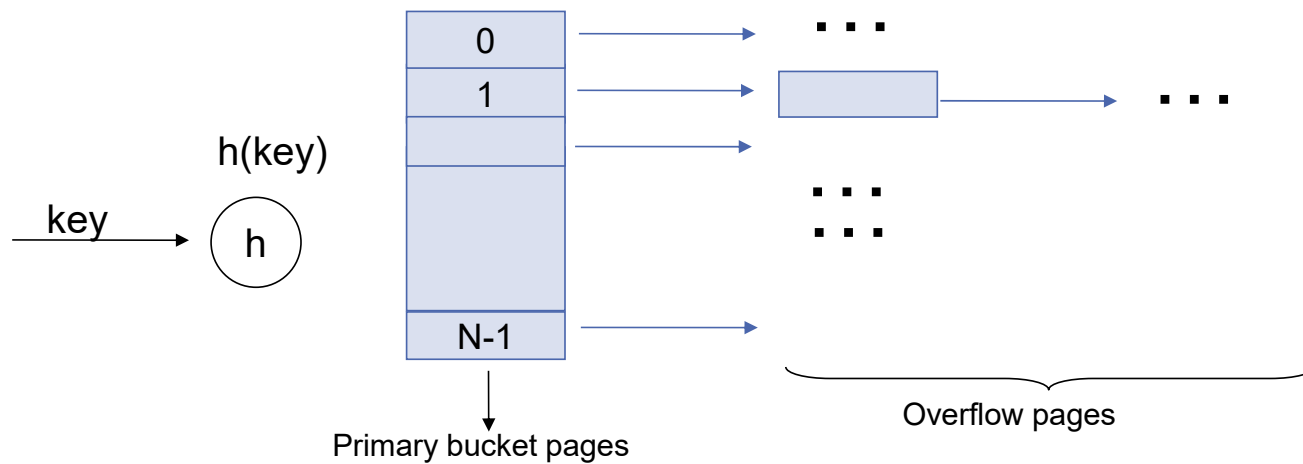
- Key density (n/T)
 - n : # of pairs in the table
 - T : Total # of possible keys
- Loading density or loading factor ($\alpha = n/(s \times b)$)
 - s : # of pairs in a bucket, b : # of buckets
- Two keys, k_1 and k_2 , are said to be *synonyms* w.r.p.t. h , if $h(k_1)=h(k_2)$.

Collisions

- Many keys might be mapped to the same home bucket
- Collision
 - When a key is mapped to a non-empty home bucket
- Overflow
 - When a key is mapped to a full home bucket
- Overflow and collision occur simultaneously when each bucket has 1 slot.

Collisions (Contd.)

- A new record hashes to a bucket that is already full
 - An overflow file is kept for storing such records
 - Overflow records that hash to each bucket can be linked together



Hashing Properties

- If # of slots is small, all operations (search, insert and delete) can be performed in $O(1)$
- Using leading letter is not a good hash function
 - Keys might bias toward certain buckets
- A good hash function should be
 - Easy to compute
 - Result few collisions

Uniform Hash Function

- A hash function that does not result in a biased use of the hash table for **random keys**
- Given a key k chosen at random, the probability that $h(k)=i$ to be $1/b$ for all buckets i
- Four popular hash functions
 - Division
 - Mid-Square
 - Folding
 - Digit Analysis

Division

- $h(k) = k \% D$
- Keys are non-negative integer
- The home bucket is obtained by using the modulo (%)
- Bucket address range from 0 to D-1
 - The hash table must have at least $b=D$ buckets
- Using a prime number for D (see textbook)

Mid-Square

■ Mid-Square:

- Squaring the keys
- Use an appropriate number of bits from the middle of the squared key as bucket address

■ If r bits is used, the size of the table is 2^r

- If there are 64 buckets (2^6), we need middle 6-bits to determine the bucket address

Folding

- The key is partitioned into several parts
- These parts are added together to obtain the key address

Digit Analysis

- All the keys in the table are known in advance
- Digitals having the most skewed distributions are deleted
- Employ the remaining digits

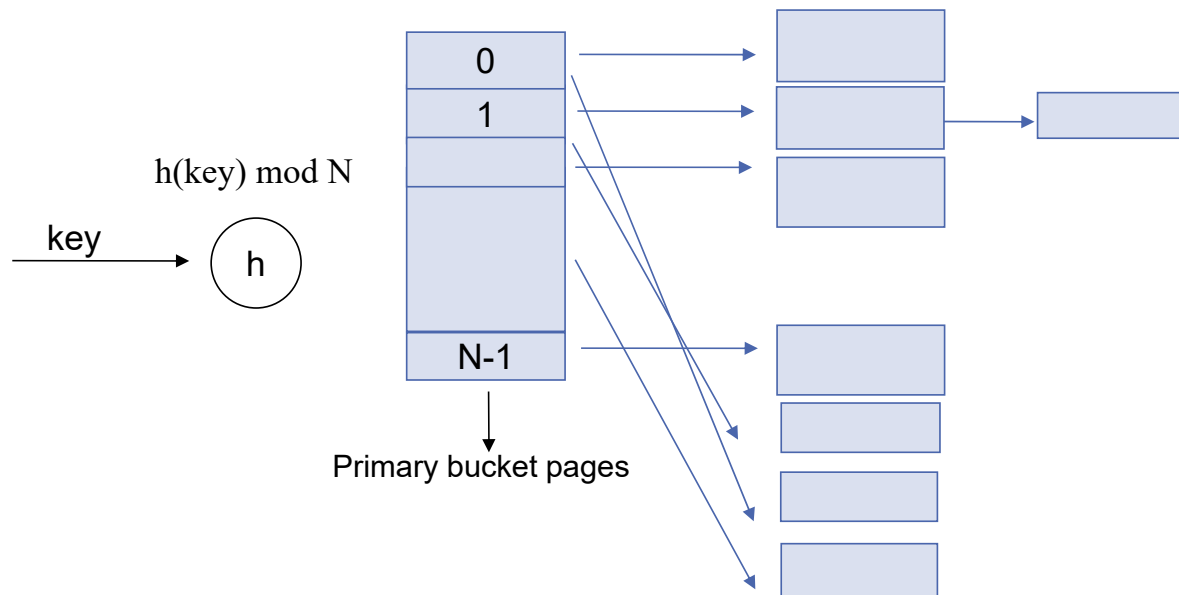


Dynamic Hashing



Dealing Overflow Problems

- Add overflow pages
- Double the size of the buckets
- Double the number of the buckets and reorganize
- ?



Dynamic Hashing

- Also called Extendable Hashing
- Use the *binary representation* of the hash value $h(K)$ in order to access a *directory*

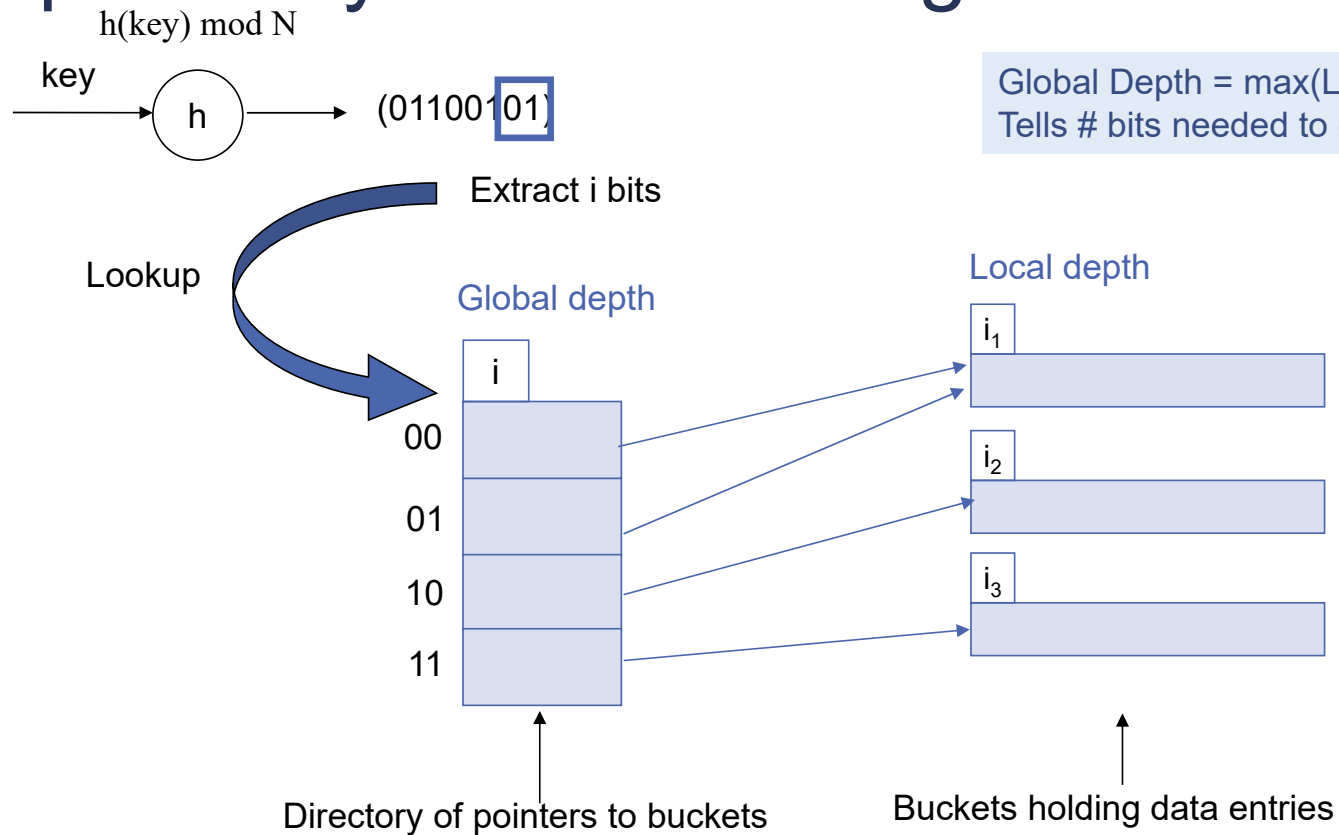
Directory

- An array of size 2^d where d is called the *global depth*
- Can be stored on disk
- Expand or shrink dynamically
- Entries point to the disk blocks
 - That contain the stored records
 - When an insertion in a disk block that is full
 - The block split into two blocks
 - The records are redistributed among the two blocks
- Updated appropriately

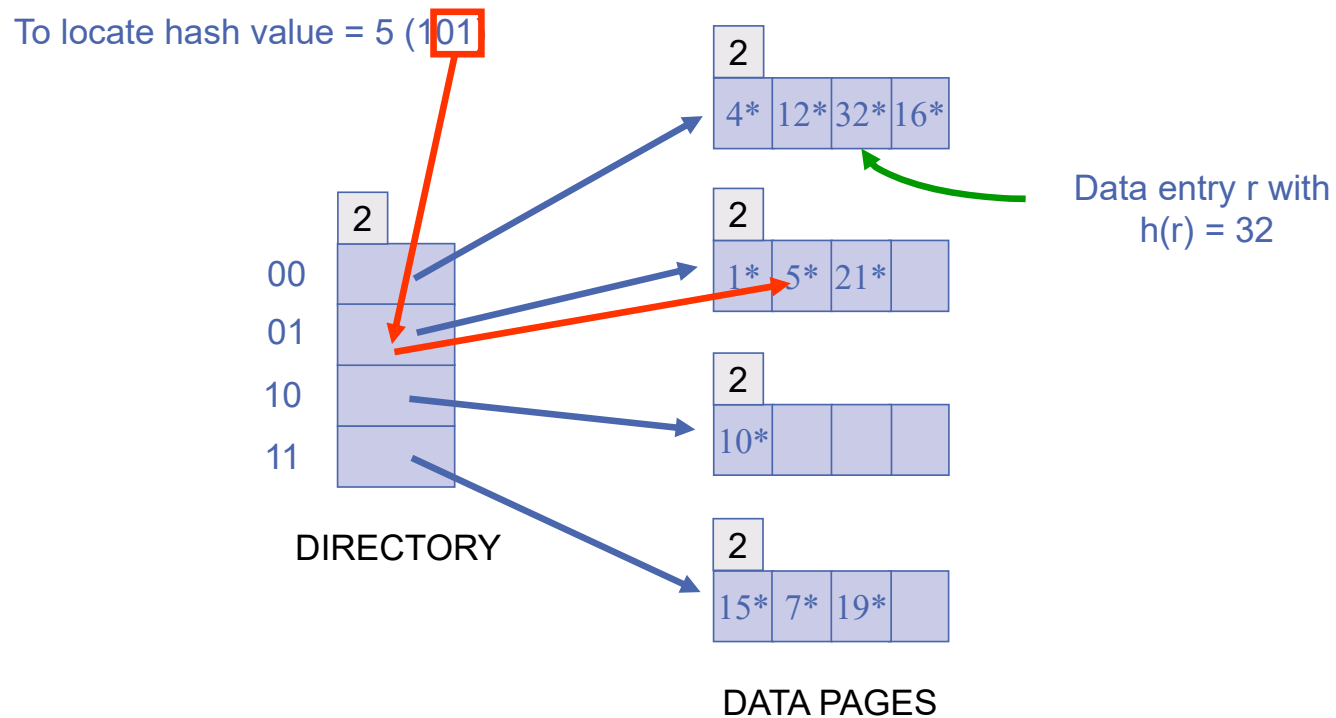
Motivation

- Situation: Bucket (primary page) becomes full. Why not re-organize file by *doubling # of buckets*?
- Idea: Use *directory of pointers to buckets*
 - Double #buckets by *doubling the directory*
 - Splitting just the bucket that overflowed!
- Directory much smaller than file
 - So doubling it is much cheaper
 - Only one page of data entries is split. *No overflow page!*

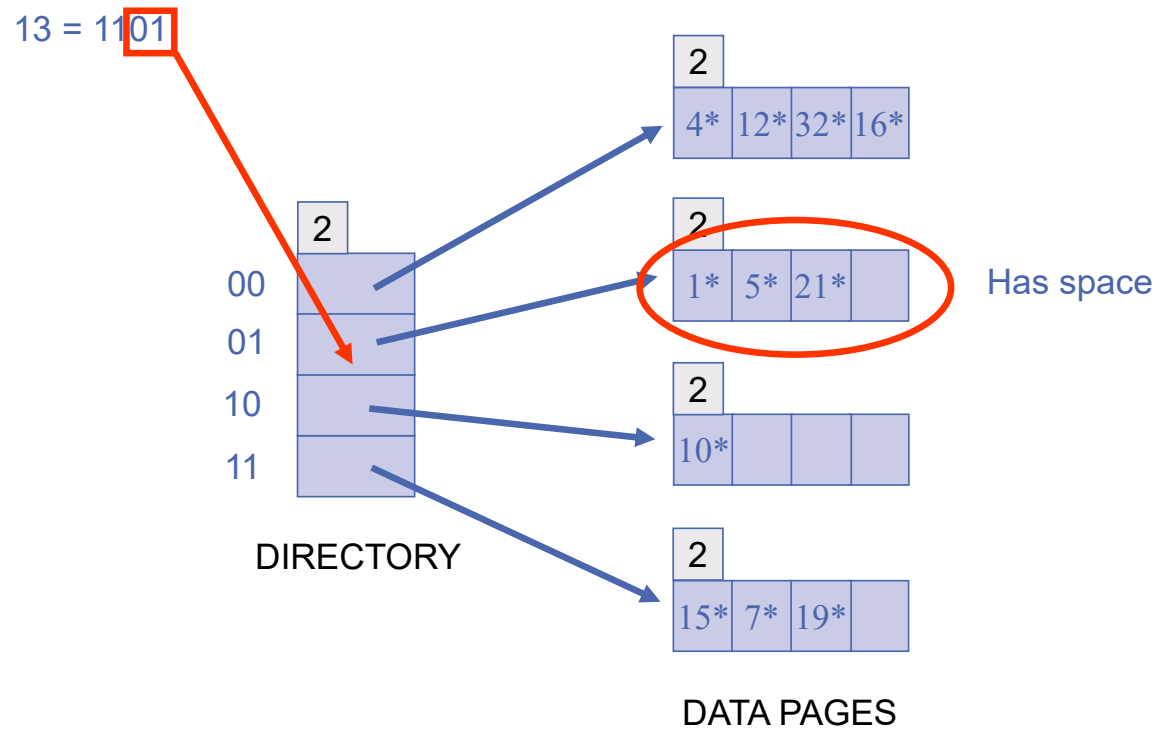
Example of Dynamic Hashing



Dynamic Hashing: Example




Dynamic Hashing: Insert 13*




Dynamic Hashing: Insert 20*

20 = 10100



Full, need split.
Consider the last
three bits

Dynamic Hashing: Insert 20*



Local/Global Depth

- Initially, all local depths are equal to global depth
 - # of bits need to express the total # of buckets
- While the process of split, if a bucket whose local depth = global depth
 - The directory must be doubled
- Global depth + 1 when the directory doubles
 - Local depth + 1 when a bucket is split




Linear Hashing



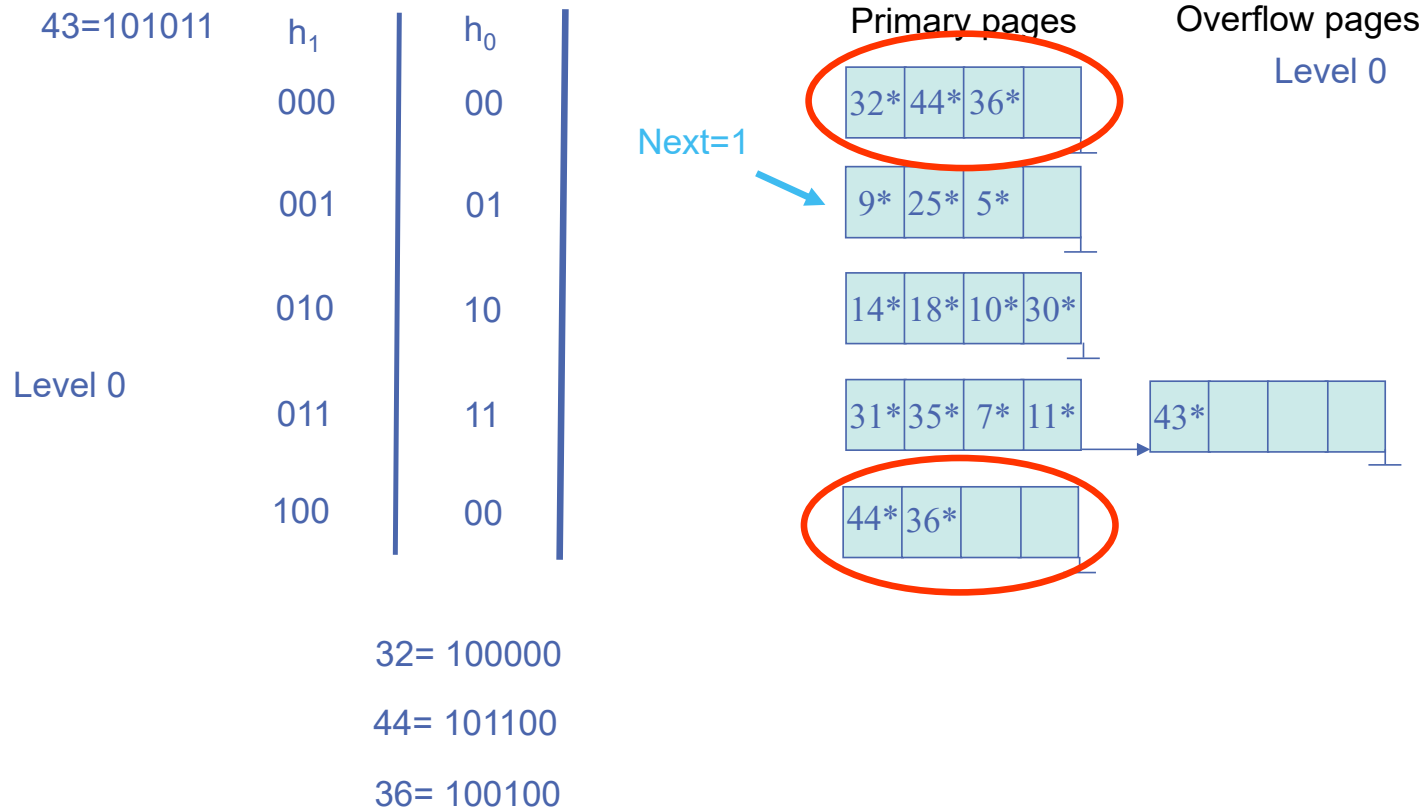
Linear Hashing

- Requires an overflow area
- Blocks are split in linear order
- Round round-robin fashion
 - eventually all buckets are split
- Idea: Use a family of hash functions h_0, h_1, h_2, \dots
 - $h_i(\text{key}) = h(\text{key}) \bmod(2^i N)$; N = initial # buckets
 - If $N = 2^{d_0}$, for some d_0 , h_i checks the last d_i bits, where $d_i = d_0 + i$.
 - h_{i+1} doubles the range of h_i

Linear Hashing: Insert 43*



Linear Hashing: Insert 43*



Linear Hashing: Insert 29*

